


Stage 1: every candidate




Da-TACOS features
(no audio)

128-d unit vector
69 s for 15,000 tracks

0.09 s for
13,000 queries

Stage 2: only the shortlist



0.2% of candidates
aligned per query

1 of 12 shifts
per pair

38 ms per query
for K = 30

$\alpha = 0.10$, $\lambda = 0.6$
(calibration works)